



CS406/CS531 Final Project

Utku Alkan
Ekin Nalbantoğlu

Algorithm 5 ParSpaRyser (*mat, cptrs, rows, cvals, dim, start, end*)

```
1:  $p \leftarrow 1$ 
2: for  $i = 1 \dots N$  do
3:    $rowSum \leftarrow 0$ 
4:   for  $j = 1 \dots N$  do
5:      $rowSum \leftarrow rowSum + mat[i, j]$ 
6:    $x[i] \leftarrow mat[i, n] - rowSum/2$ 
7:    $p \leftarrow p \times x[i]$ 
8: for each thread do
9:    $myX \leftarrow x$ 
10:   $myP \leftarrow 0$ 
11:   $myStart \leftarrow start + threadId \times chunkSize$ 
12:   $myEnd \leftarrow \min(start + (threadId + 1) \times chunkSize, end)$ 
13:  calculate  $myX$  using  $GrayCode_{myStart-1}$ 
14:   $prod \leftarrow \prod_{n=1}^{dim} myX[i]$  for  $myX[i] \neq 0$ 
15:   $zeroNum \leftarrow \sum_{n=1}^{length(myX)} myX[i]$  for  $myX[i] = 0$ 
16:  for  $g = myStart \dots myEnd - 1$  do
17:     $j \leftarrow \log_2(GrayCode_g \oplus GrayCode_{g-1}) + 1$ 
18:     $s \leftarrow 2 * GrayCode_g[j] - 1$ 
19:    for  $i = cptrs[j] \dots cptrs[j]$  do
20:      if  $myX[rows[i]] == 0$  then
21:         $zeroNum \leftarrow zeroNum - 1$ 
22:         $myX[rows[i]] \leftarrow myX[rows[i]] + s \times cvals[i]$ 
23:         $prod \leftarrow prod \times myX[rows[i]]$ 
24:      else
25:         $prod \leftarrow prod / myX[rows[i]]$ 
26:         $myX[rows[i]] \leftarrow myX[rows[i]] + s \times cvals[i]$ 
27:        if  $myX[rows[i]] == 0$  then
28:           $zeroNum \leftarrow zeroNum + 1$ 
29:        else
30:           $prod \leftarrow prod \times myX[rows[i]]$ 
31:      if  $zeroNum == 0$  then
32:         $myP \leftarrow myP + (-1)^g * prod$ 
33:       $AtomicAdd(p, myP)$ 
34: return  $p \times (4 \times (n \bmod 2) - 2)$ 
```

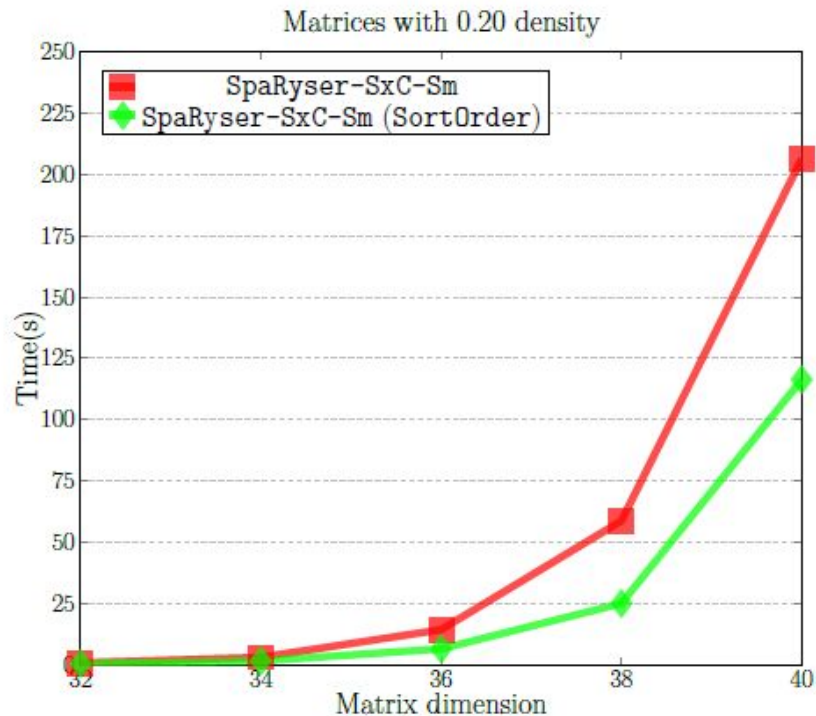

COMPUTING MATRIX PERMANENTS AND COUNTING PERFECT MATCHINGS ON GPUS


by
B. Yağlıoğlu

Algorithm 5 ParSpaRyser: Ryser
algorithm to calculate
permanents of sparse matrices
with a parallelizable block.

Preprocessing

```
sortMatrixColumnsByNonZeros(matrix, n);
```





Integration of dynamic scheduling by dividing the load into chunks


```
502     long long chunkSize = (end - start) / numChunks + 1;
503
504     #pragma omp parallel num_threads(numThreads)
505     {
506         double myX[N];
507         #pragma omp for schedule(dynamic, 1)
508         for (int chunkID = 0; chunkID < numChunks; chunkID++)
509         {
510
511             for (int i = 0; i < N; i++) {
```

Determining the chunk size

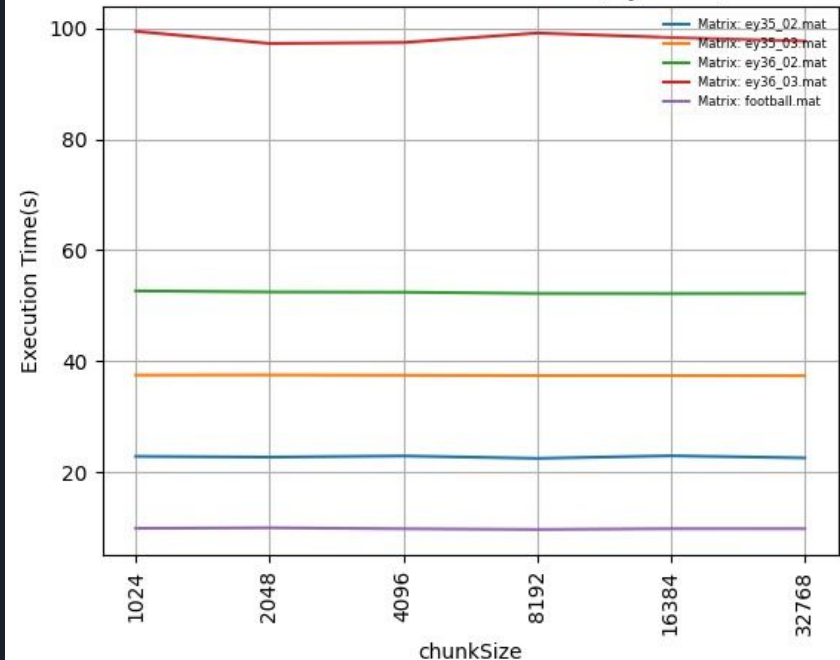
```
numChunks = 8192;
```

```
start = omp_get_wtime();  
double perm_spaRyser = ParSpaRyserD(ccs.cptrs,  
end = omp_get_wtime());  
cout << perm_spaRyser << "\t" << end-start << endl;
```

Execution Time vs Dimensions (Guided)




Execution Time vs Dimensions (Dynamic)



CUDA implementation kernel call for 1 gpu

```
558     int gridsize = 16384;  
559     int blocksize = 256;
```

```
long long int start = 1;  
long long int end = (1LL << (N-1));  
int numThreads = numBlocks * blockSize;  
long long int chunkSize = (end - start) / numThreads + 1;  
  
ParSpaRyserCudaKernel<<<numBlocks, blockSize>>>(d_cptrs, d_rows,  
d_cvals, N, d_p, d_x, 1, (1LL << (N-1)), chunkSize);  
  
cudaDeviceSynchronize();
```



2 and 4 GPUs

```
if(gpuCounter == 0){
    ParSpaRyserCudaKernel<<<numBlocks, blockSize>>>(d_cptrs, d_rows,
    d_cvals, N, d_p, d_x, start, (end/2), chunkSize);
}else{
    ParSpaRyserCudaKernel<<<numBlocks, blockSize>>>(d_cptrs, d_rows,
    d_cvals, N, d_p, d_x, end/2, end, chunkSize);
}
```

```
double startTime = omp_get_wtime();
if(gpuCounter == 0){
    ParSpaRyserCudaKernel<<<numBlocks, blockSize>>>(d_cptrs,
    d_rows, d_cvals, N, d_p, d_x, start, 3*end/8, chunkSize);
}else if(gpuCounter==1){
    ParSpaRyserCudaKernel<<<numBlocks, blockSize>>>(d_cptrs,
    d_rows, d_cvals, N, d_p, d_x, 3*end/8, 6*end/8, chunkSize);
}else if(gpuCounter==2){
    ParSpaRyserCudaKernel<<<numBlocks, blockSize>>>(d_cptrs,
    d_rows, d_cvals, N, d_p, d_x, 6*end/8, 7*end/8, chunkSize);
}else if(gpuCounter==3){
    ParSpaRyserCudaKernel<<<numBlocks, blockSize>>>(d_cptrs,
    d_rows, d_cvals, N, d_p, d_x, 7*end/8, end, chunkSize);
}
```

```
cs406.ekinn@nebula:~/proj/twoGPU$ make
nvcc -Xcompiler -fopenmp -o gpu4 parSpaRyserFourGPU.
cu
./gpu4 chesapeake.mat
GPU ID: 3 synchronizes in 16.781465 seconds
GPU ID: 2 synchronizes in 17.602388 seconds
GPU ID: 1 synchronizes in 19.010036 seconds
GPU ID: 0 synchronizes in 19.458299 seconds
spaRyser Result: 1.31735e+13
20.318988 seconds
```



Thank you

Kaya, K. (2019). Parallel algorithms for computing sparse matrix permanents. Turkish Journal of Electrical Engineering and Computer Sciences, 27, 4284–4297.

Yağlıoğlu, B. (2021). Computing matrix permanents and counting perfect matchings on gpus [Master's thesis]. <https://research.sabanciuniv.edu/id/eprint/42490/1/10364196.pdf>